

Perfect question 🙌 — understanding the **general structure** of a Kubernetes YAML file is the foundation for creating any manifest (Deployment, Service, ConfigMap, etc.).

Let's go step by step 👇

✳️ General Structure of a Kubernetes YAML File

Every Kubernetes YAML file follows this **four-part structure**:

```
apiVersion: <API version>  
kind: <Resource type>  
metadata:  
  <information about the object>  
spec:  
  <desired configuration / behavior>
```

🧱 1 apiVersion

Defines **which Kubernetes API version** the resource belongs to.

Common examples:

Resource	apiVersion
Pod	v1
Service	v1
Deployment	apps/v1
ConfigMap	v1
Ingress	networking.k8s.io/v1

📘 Think of this as "**which K8s API endpoint handles this file**".

🧱 2 kind

Defines **what type of resource** you're creating.

Examples:

```
Pod
```

Deployment

Service

ConfigMap

Secret

Ingress

Namespace

 It tells Kubernetes **what object** you want to create.

3 metadata

Holds identifying information like the object's **name**, **labels**, and **namespace**.

Example:

```
metadata:
```

```
  name: myapp
```

```
  namespace: default
```

```
  labels:
```

```
    app: myapp
```

 Think of it as the resource's **identity card**.

spec

Describes the **desired state** or configuration details for the resource.

What goes inside spec: depends on the resource type (kind).

Example 1: Deployment (full example)

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

```
  name: nodeapp-deployment
```

```
  labels:
```

```
app: nodeapp
```

```
spec:
```

```
replicas: 2
```

```
selector:
```

```
matchLabels:
```

```
    app: nodeapp
```

```
template:
```

```
metadata:
```

```
labels:
```

```
    app: nodeapp
```

```
spec:
```

```
containers:
```

```
- name: nodeapp
```

```
    image: nodeapp:v1
```

```
ports:
```

```
- containerPort: 3000
```

✳ Example 2: Service

```
apiVersion: v1
```

```
kind: Service
```

```
metadata:
```

```
    name: nodeapp-service
```

```
spec:
```

```
    type: NodePort
```

```
    selector:
```

```
        app: nodeapp
```

```
    ports:
```

```
        - port: 3000
```

```
targetPort: 3000
```

```
nodePort: 30080
```

🧠 Optional Fields (Appear Sometimes)

Field	Purpose	Example
status	Shows current state (added by K8s, not by you)	status: Running
data	For storing config values (in ConfigMaps, Secrets)	data: { KEY: VALUE }
rules	For RBAC, NetworkPolicy, Ingress	Defines access rules

⚡ Summary Structure (Template)

Here's a general reusable structure you can follow:

```
apiVersion: <api-version>  
kind: <resource-kind>  
metadata:  
  name: <resource-name>  
  namespace: <namespace-name>  # optional  
  labels:  
    key: value  
spec:  
  <resource-specific configuration>
```

⚡ Example in One Line Explanation

“Every Kubernetes YAML file defines **what you want (spec)**, **what it is (kind)**, **where it belongs (metadata)**, and **who handles it (apiVersion)**.”

Perfect 🔥 — here's your **Kubernetes YAML Cheat Sheet** — a reusable **template library** for all the major K8s objects you'll use daily as a DevOps or developer.

Each block includes the **structure**, **purpose**, and a **ready-to-use example** ✅

✳️ 1 Pod

A Pod is the smallest deployable unit in Kubernetes — runs one or more containers.

```
apiVersion: v1
kind: Pod
metadata:
  name: mypod
  labels:
    app: myapp
spec:
  containers:
  - name: mycontainer
    image: nginx:latest
    ports:
    - containerPort: 80
```

✳️ 2 Deployment

Ensures a desired number of identical Pods (replicas) are always running.
Handles rolling updates and rollbacks.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: myapp-deployment
  labels:
    app: myapp
```

```
spec:  
  replicas: 3  
  selector:  
    matchLabels:  
      app: myapp  
  template:  
    metadata:  
      labels:  
        app: myapp  
    spec:  
      containers:  
        - name: myapp  
          image: nginx:1.25  
          ports:  
            - containerPort: 80  
          livenessProbe:  
            httpGet:  
              path: /  
              port: 80  
            initialDelaySeconds: 5  
            periodSeconds: 10
```

3 ReplicaSet

Ensures a fixed number of identical Pods are running.
Usually managed automatically by a Deployment.

```
apiVersion: apps/v1  
kind: ReplicaSet  
metadata:
```

```
name: myapp-replicaset

spec:

replicas: 2

selector:

matchLabels:

  app: myapp

template:

metadata:

labels:

  app: myapp

spec:

containers:

- name: myapp

  image: nginx:latest
```

Service

Exposes your Pods (internally or externally) so other apps/users can access them.

Common Types:

- **ClusterIP (default):** Internal only
- **NodePort:** Accessible from outside via <NodeIP>:<Port>
- **LoadBalancer:** Cloud-based external access

```
apiVersion: v1

kind: Service

metadata:

  name: myapp-service

spec:

  type: NodePort

  selector:
```

```
app: myapp

ports:
- port: 80
  targetPort: 80
  nodePort: 30080
```

5 ConfigMap

Stores **non-sensitive configuration data** as key-value pairs.
Used for environment variables or config files.

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: myapp-config
data:
  APP_ENV: "production"
  APP_VERSION: "v1.0"
```

Attach it to a Pod/Deployment:

```
envFrom:
- configMapRef:
    name: myapp-config
```

6 Secret

Stores **sensitive information** like passwords, API keys, etc.
Values must be base64-encoded.

```
apiVersion: v1
kind: Secret
metadata:
  name: myapp-secret
```

```
type: Opaque  
  
data:  
  
DB_USER: dXNlcg== # base64 for 'user'  
  
DB_PASS: cGFzc3dvcmQ= # base64 for 'password'  
  
Use it in a Deployment:  
  
envFrom:  
  
- secretRef:  
  
  name: myapp-secret
```

7 Ingress

Provides **domain-based routing** (e.g., myapp.local → Service).
Requires an Ingress Controller (e.g., NGINX).

```
apiVersion: networking.k8s.io/v1  
  
kind: Ingress  
  
metadata:  
  
  name: myapp-ingress  
  
spec:  
  
  rules:  
  
  - host: myapp.local  
  
    http:  
  
      paths:  
  
      - path: /  
  
        pathType: Prefix  
  
      backend:  
  
        service:  
  
          name: myapp-service  
  
        port:  
  
          number: 80
```

8 Namespace

Logical grouping for separating environments (dev, staging, prod).

```
apiVersion: v1
kind: Namespace
metadata:
  name: dev-environment
```

9 PersistentVolumeClaim (PVC)

Requests storage from the cluster (used by databases or file storage apps).

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: myapp-pvc
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 1Gi
  Use in Pod:
  volumes:
    - name: data-storage
      persistentVolumeClaim:
        claimName: myapp-pvc
```

10 Job

Runs a **one-time task** (like batch processing or database migration).

```
apiVersion: batch/v1
kind: Job
metadata:
  name: data-migration-job
spec:
  template:
    spec:
      containers:
        - name: migration
          image: busybox
          command: ["echo", "Running data migration"]
      restartPolicy: Never
      backoffLimit: 2
```

Quick Recap: YAML Skeleton

```
apiVersion: <api-version>
kind: <resource-kind>
metadata:
  name: <resource-name>
  labels:
    key: value
spec:
  <resource-specific-configuration>
```

Tip: Apply All at Once

If you have multiple YAML files (deployment, service, configmap, etc.), you can apply them together:

```
kubectl apply -f .
```

Or a single combined file separated by ---:

```
# deployment.yaml
apiVersion: apps/v1
kind: Deployment
...
---
# service.yaml
apiVersion: v1
kind: Service
...
```

Excellent 🌟 — this is one of the most important concepts in Kubernetes.
Let's break it down **in very simple language** first, then we'll make a clear **summary table** at the end 🤝

✳️ What Are Kubernetes Objects?

Kubernetes objects are like **building blocks** — each object represents **something that exists in your cluster**.

👉 You can think of them as **records or instructions** that tell Kubernetes what you want your system to look like.

Each object:

- Describes a **desired state** (what you want).
 - Kubernetes **constantly works** to make the cluster match that state.
-

🧠 The Main Kubernetes Objects (Explained Simply)

1 Pod

- The smallest unit in Kubernetes.
- It runs **one or more containers** (like a mini environment for your app).
- Example: A Pod running your Node.js app.

📦 *Think of it like a box that holds your app container.*

2 ReplicaSet

- Makes sure a **specific number of Pods** are always running.
- If one Pod crashes, it **creates a new one** automatically.

⌚ *Think of it like a backup manager that ensures the right number of Pods.*

3 Deployment

- A higher-level controller that **manages ReplicaSets**.
- Helps you **update your app smoothly** (rolling updates, rollbacks).
- You mostly use Deployment instead of ReplicaSet directly.

 *Think of it as your app's manager that handles scaling and version updates.*

4 Service

- Provides a **stable way to access Pods**, even if Pods keep changing.
- It has a **fixed IP** and can **load-balance traffic** between multiple Pods.

 *Think of it like a receptionist that forwards users to available Pods.*

5 ConfigMap

- Stores **non-secret configuration data** (like environment variables, URLs).
- Keeps your app config separate from the app code.

 *Think of it as a notepad that stores app settings.*

6 Secret

- Stores **sensitive data** (like passwords or API keys) securely.
- Values are encoded (base64).

 *Think of it as a locked box for private info.*

7 Namespace

- Divides your cluster into **separate sections** for organization.
- Useful for separating dev, test, and prod environments.

 *Think of it as folders in your computer.*

8 Ingress

- Manages **external HTTP/HTTPS access** to your services.

- Helps with **domain-based routing** (e.g., myapp.com → app-service).

 Think of it as your app's front gate or entry point.

PersistentVolume (PV) and PersistentVolumeClaim (PVC)

- Used for **storing data** permanently (like databases).
- PV = actual storage.
PVC = request for that storage.

 Think of PV as a hard drive and PVC as asking to use part of that drive.

Job / CronJob

- **Job:** Runs a task once (like a script or backup).
- **CronJob:** Runs it on a schedule (like every night).

 Think of it like a scheduled task on your computer.

Bonus: Controller vs Object

- **Objects** = the *things that exist* (Pods, Services, etc.)
 - **Controllers** = the *brains* that make sure objects match your desired state (Deployment, ReplicaSet, etc.)
-

Summary Table of Kubernetes Objects

Object	Purpose	Example Use	Simple Analogy
Pod	Runs containers	Node.js app container	Box holding your app
ReplicaSet	Ensures fixed number of Pods	Always 3 Pods running	Backup manager
Deployment	Manages Pods/ReplicaSets, handles updates	Deploying new app version	Project manager

Service	Provides stable network access	Load balance between Pods	Reception desk
ConfigMap	Stores configuration data	Database URL, environment	Notepad for settings
Secret	Stores sensitive data	API key, password	Locked box
Namespace	Divides cluster logically	dev, staging, prod	Folder system
Ingress	Manages external access (HTTP/HTTPS)	Route myapp.com → service	Front gate
PersistentVolume (PV)	Physical storage	Disk in cluster	Hard drive
PersistentVolumeClaim (PVC)	Requests storage	App data request	Borrow part of hard drive
Job	One-time task	Data migration	Manual task
CronJob	Repeated/scheduled task	Daily backups	Alarm clock task

 **In One Line:**

Kubernetes objects are the **blueprints** of your cluster — you declare what you want, and Kubernetes automatically builds and maintains it.

Would you like me to show you a **visual hierarchy diagram (Pods → ReplicaSet → Deployment → Service → Ingress)** to understand how these objects connect together?

Object	Purpose	Who Manages It	Example Use
Pod	Runs one or more containers	Deployment / ReplicaSet	Run your Node.js app
ReplicaSet	Keeps the right number of Pods alive	Deployment	Maintain 3 running copies
Deployment	Manages app rollout & versioning	You (user-defined)	Update from v1 to v2
Service	Provides stable access to Pods	Kubernetes	Expose app on cluster IP
Ingress	Routes external HTTP traffic	Kubernetes	myapp.com → backend service
ConfigMap	Stores non-secret configs	Deployment mounts it	Store environment variables
Secret	Stores sensitive info	Deployment mounts it	Store DB password
Namespace	Groups related objects	You / Admin	Dev, Test, Prod isolation
PersistentVolume (PV)	Actual storage resource	Admin	Disk or cloud volume
PersistentVolumeClaim (PVC)	Request for PV storage	App / User	"I need 5GB for uploads"